

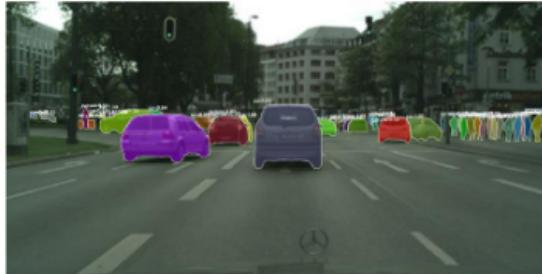
Semantic Segmentation and Instance Segmentation

Mask R-CNN

By Kaiming He, Georgia Gkioxari, Piotr Dollar and Ross Girshick
Presented By Aditya Sanghi

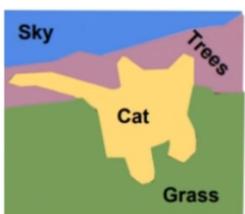
<https://www.cs.toronto.edu/~fidler/teaching/2018/slides/CSC2548/MaskRCNN.pdf>

Semantic segmentation --- > instance segmentation



Object
Detection

Instance
Segmentation



This image is CC0 public domain

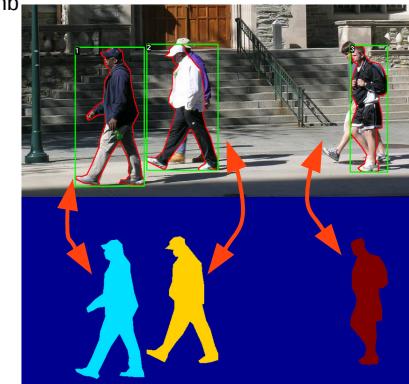
DOG, DOG, CAT



DOG, DOG, CAT

Multiple Object

https://colab.research.google.com/github/pytorch/tutorials/blob/gh-pages/_downloads/torchvision_finetuning_instance_segmentation.ipynb



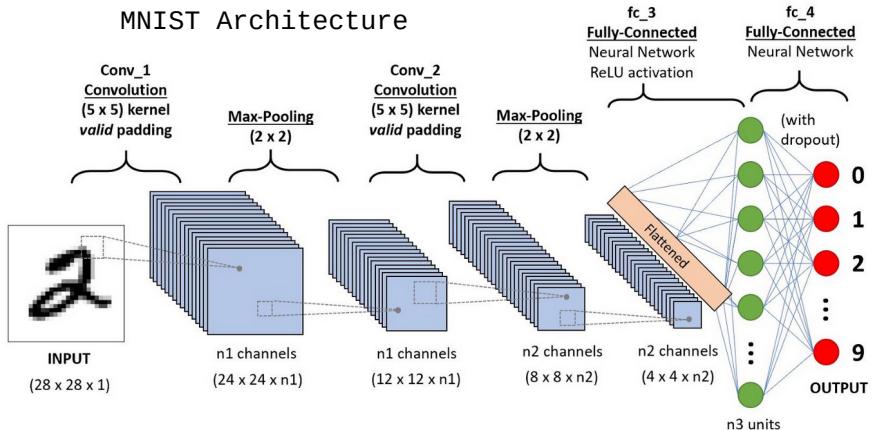
segmentation masks

Each image has a corresponding segmentation mask, where each color correspond to a different instance.

Reference: Detecting Pedestrians using PyTorch – A Helpful Guide

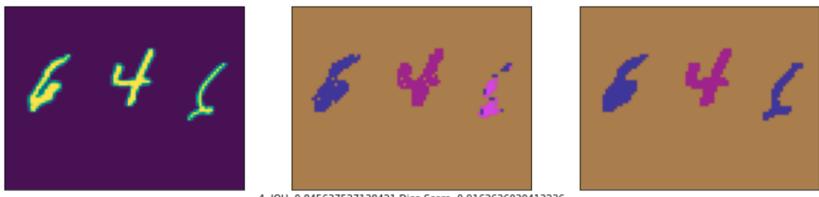
<https://www.data-blogger.com/pedestrian-detection-using-pytorch/>

From MNIST to Semantic Segmentation of Handwritten Digits



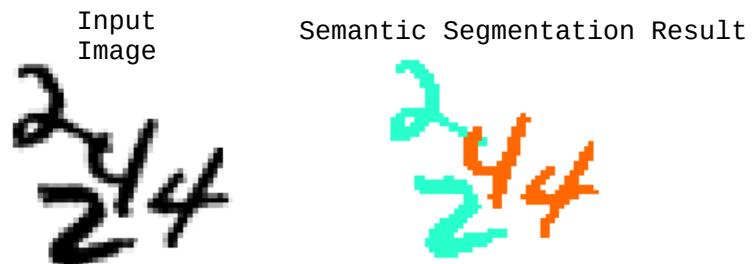
sinarazi /
Image-Segmentation-of-Handwritten-Di
gits
<https://github.com/sinarazi/Image-Segmentation-of-Handwritten-Digits>

Semantic segmentation: to predict/classify the handwritten digits with segmentation mask on each pixel basis



MNIST extended: a dataset for semantic segmentation

<https://awaywithideas.com/mnist-extended-a-dataset-for-semantic-segmentation-and-object-detection/>



Four Step Design Guidelines

General Guide (4 Steps)

1. Replace FC layers with convolutional layers.
2. Convert the last layer output to the original resolution.
3. Do softmax-cross entropy between the pixelwise predictions and segmentation ground truth.
4. Backprop and SGD

Stochastic Gradient Descent (SGD)

Entropy:

$$H(X) = - \sum_{i=1}^n P(x_i) \log P(x_i) \quad (1)$$

Cross Entropy: Measure of difference between 2 probabilities

$$H(p, q) = - \sum_{x \in \mathcal{X}} p(x) \log q(x) \quad (2)$$

Define loss function (Cross entropy)

$$D(\hat{Y}, Y) = - \frac{1}{N} \cdot \sum Y_i \cdot \log(\hat{Y}_i) \quad (3)$$

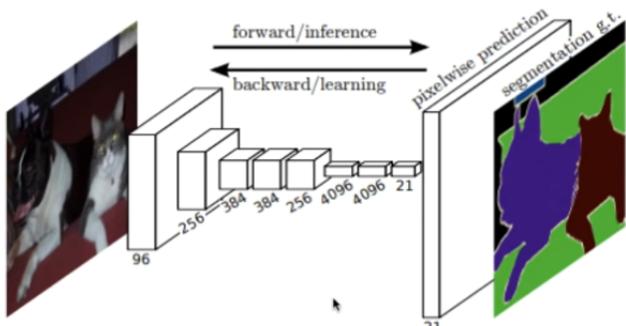
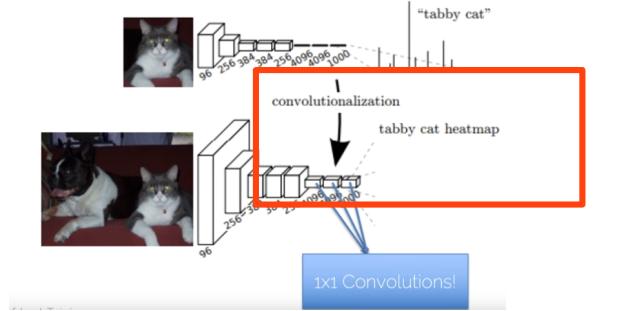
$$Q(w) = \sum_{i=1}^n Q_i(w) \quad (4) \text{ Objective function}$$

$$w := w - \eta \nabla Q(w) = w - \frac{\eta}{n} \sum_{i=1}^n \nabla Q_i(w) \quad (5) \text{ SGD}$$

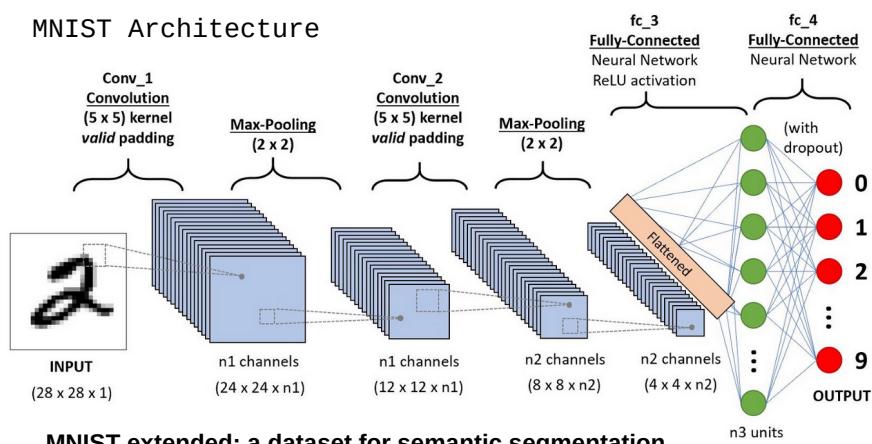
Convolutionalization: Change FC Layer to Convolutional Layer

General Guide: change FC (fully connected) layer to convolutional layer by 1×1 convolution kernel

"Convolutionalization"



MNIST Architecture



MNIST extended: a dataset for semantic segmentation

<https://awaywithideas.com/mnist-extended-a-dataset-for-semantic-segmentation-and-object-detection/>

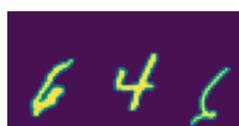
Input Image



Semantic Segmentation Result



Semantic segmentation: to predict/classify the hand written digits with segmentation mask on each pixel basis



SGD Example

Example [\[edit\]](#)

Let's suppose we want to fit a straight line $\hat{y} = w_1 + w_2 x$ to a training set with observations (x_1, x_2, \dots, x_n) and corresponding estimated responses $(\hat{y}_1, \hat{y}_2, \dots, \hat{y}_n)$ using [least squares](#). The objective function to be minimized is:

$$Q(w) = \sum_{i=1}^n Q_i(w) = \sum_{i=1}^n (\hat{y}_i - y_i)^2 = \sum_{i=1}^n (w_1 + w_2 x_i - y_i)^2.$$

The last line in the above pseudocode for this specific problem will become:

$$\begin{bmatrix} w_1 \\ w_2 \end{bmatrix} := \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} - \eta \begin{bmatrix} \frac{\partial}{\partial w_1} (w_1 + w_2 x_i - y_i)^2 \\ \frac{\partial}{\partial w_2} (w_1 + w_2 x_i - y_i)^2 \end{bmatrix} = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} - \eta \begin{bmatrix} 2(w_1 + w_2 x_i - y_i) \\ 2x_i(w_1 + w_2 x_i - y_i) \end{bmatrix}.$$

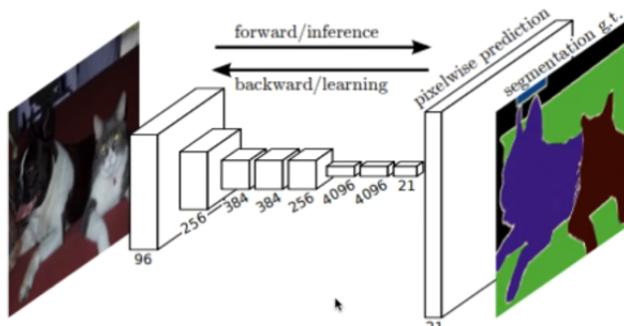
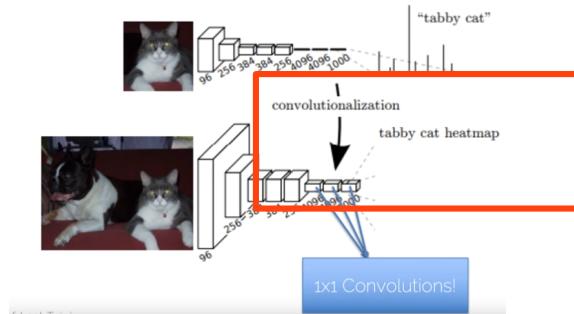
Note that in each iteration (also called update), the gradient is only evaluated at a single point x_i instead of at the set of all samples.

The key difference compared to standard (Batch) Gradient Descent is that only one piece of data from the dataset is used to calculate the step, and the piece of data is picked randomly at each step.

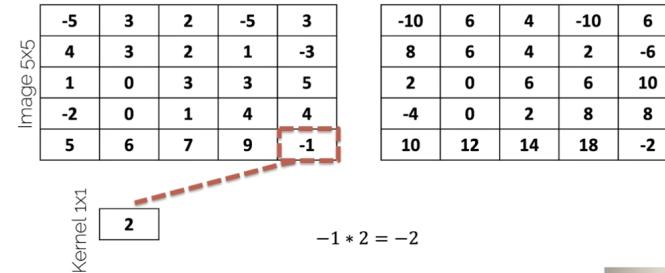
Convolutionalization: Change FC Layer to Convolutional Layer

General Guide: change FC (fully connected) layer to convolutional layer by 1x1 convolution kernel

"Convolutionalization"



1x1 Convolution



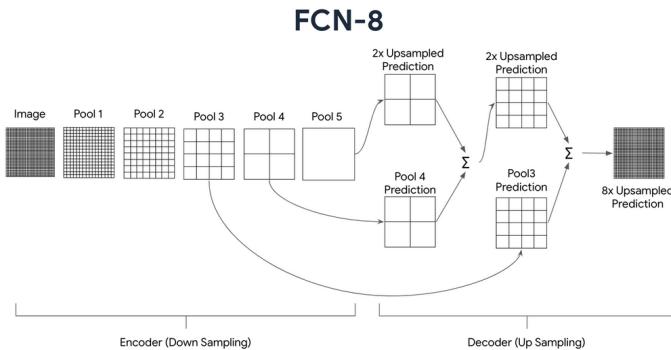
1x1 Convolution

[Li et al. 2013]

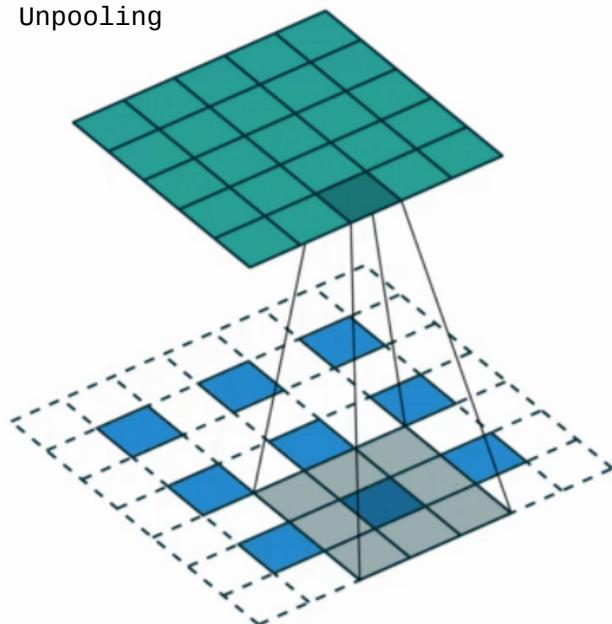


Network in Network

Up Sampling



Unpooling



- 1. Interpolation

Original image x 10

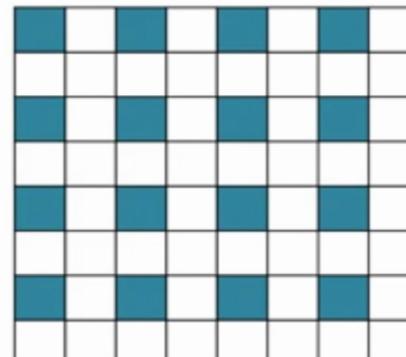
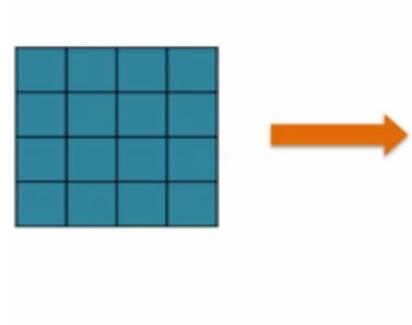


Nearest neighbor interpolation

Bilinear interpolation

Bicubic inter

- 2. Fixed unpooling



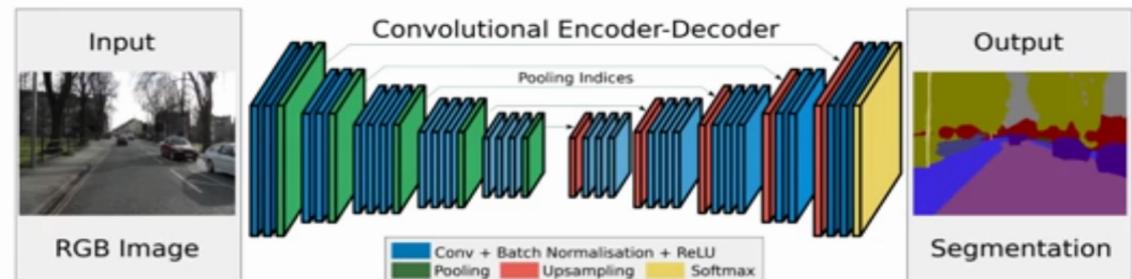
RCNN Architecture Examples

<https://www.youtube.com/watch?v=XMSjOatyH0k>

Mask R-CNN is a state of the art model for instance segmentation, developed on top of Faster R-CNN. Faster R-CNN is a region-based convolutional neural networks [2], that returns bounding boxes for each object and its class label with a confidence score.

1. Using [Mask R-CNN](#), which is based on top of [Faster R-CNN](#). Faster R-CNN is a model that predicts both bounding boxes and class scores for potential objects in the image.
2. Mask R-CNN adds an extra branch into Faster R-CNN, which also predicts segmentation masks for each instance.

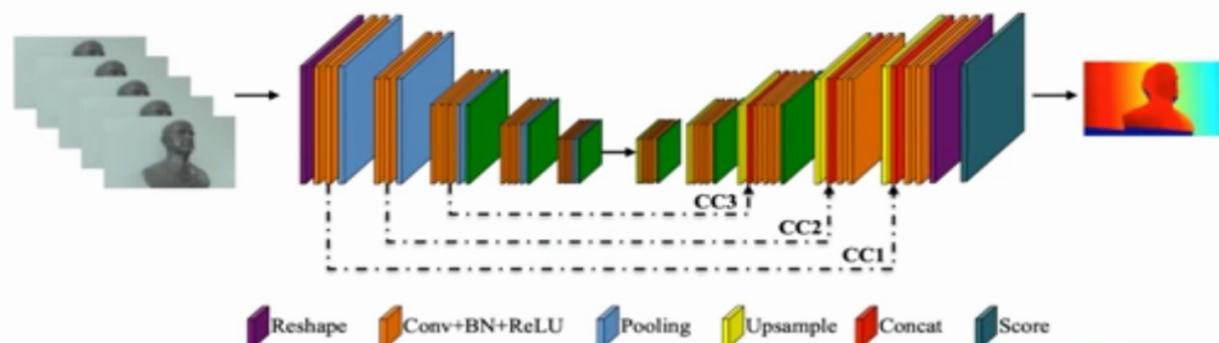
Architecture: Deep Depth from Focus.



Badrinarayanan et al., SegNet: A Deep Convolutional Encoder-Decoder Architecture for Image Segmentation

Skip Connections

- Concatenation connections



Mask R-CNN for Semantic Segmentation

<https://github.com/buseyaren/Installation-MaskRCNN>

Step 6: Running the setup.py file. \$python3 setup.py installpip

Step 7. install git+

<https://github.com/phiferriere/cocoapi.git#subdirectory=Python API>

Step 8. \$python3 balloon.py train --dataset=/path/to/dataset --model=coco

Check tensorflow in Python before down loading:

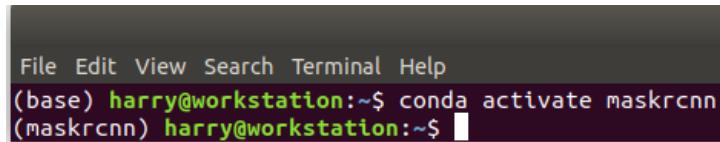
```
import tensorflow as tf  
print(tf.__version__)
```

<https://www.youtube.com/watch?v=tcu4pr948n0>

Software architecture: <https://blog.paperspace.com/mask-r-cnn-in-tensorflow-2-0/>

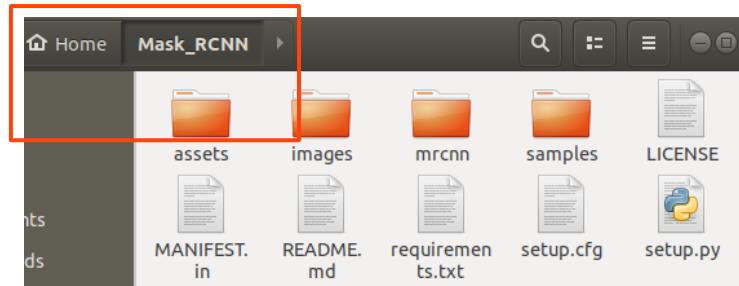
Step 1. use conda to create a virtual environment \$conda create -n maskrcnn python=3.6.12

Step 2. Activate the environment: \$conda activate maskrcnn



```
File Edit View Search Terminal Help  
(base) harry@workstation:~$ conda activate maskrcnn  
(maskrcnn) harry@workstation:~$
```

Step 3. \$git clone https://github.com/matterport/Mask_RCNN.git



Step 4. python3 setup.py install

Step 4. \$pip3 install -r requirements.txt (add the following 2 lines to enable gpu: (add this line) tensorflow-gpu==1.15.0 (then add this line) keras==2.2.5

Step 5. Download the pre-trained weights from https://github.com/matterport/Mask_RCNN/releases. Download the file mask_rcnn_balloon.h5 from Mask_RCNN_2.1 file. These 2 models should be placed in the samples folder.
download: [mask_rcnn_coco.h5](#)

Pedestrian Detection

<https://data-flair.training/blogs/pedestrian-detection-python-opencv/>

TorchVision Instance Segmentation

Step 1. start colab, install pycocotools-2.0
COCO metric for intersection over union.

```
%shell
pip install cython
# Install pycocotools, the version by default in
# has a bug fixed in https://github.com/cocodatas
pip install -U 'git+https://github.com/cocodatas'
```

Cursor over it to run the code above on colab

dataset for Penn-Fudan, download and extract the data

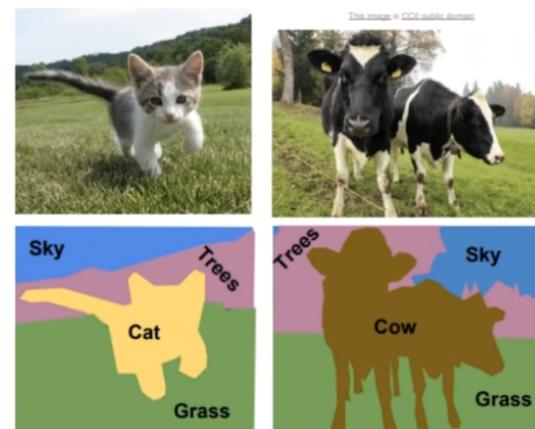
https://www.cis.upenn.edu/~jshi/ped_html/PennFudanPed.zip

The data is structured as follows

```
PennFudanPed/
  PedMasks/
    FudanPed00001_mask.png
    FudanPed00002_mask.png
    FudanPed00003_mask.png
    FudanPed00004_mask.png
    ...
  PNGImages/
    FudanPed00001.png
    FudanPed00002.png
    FudanPed00003.png
```

1. Using [Mask R-CNN](#), which is based on top of [Faster R-CNN](#). Faster R-CNN is a model that predicts both bounding boxes and class scores for potential objects in the image.
2. Mask R-CNN adds an extra branch into Faster R-CNN, which also predicts segmentation masks for each instance.

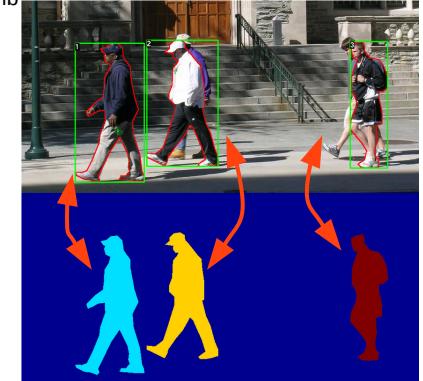
Semantic segmentation --- > instance segmentation



Detecting Pedestrians using PyTorch – A Helpful Guide

<https://www.data-blogger.com/pedestrian-detection-using-pytorch/>

https://colab.research.google.com/github/pytorch/tutorials/blob/gh-pages/_downloads/torchvision_finetuning_instance_segmentation.ipynb



segmentation masks

Each image has a corresponding segmentation mask, where each color correspond to a different instance.